Codex session 01a18f11…

• Project: Sonrisa

• Storage: active

• Session ID: 01a18f11-f52d-7401-850a-b62d1de60fad

• Started: 2026-10-30T23:08:01.291Z

• Updated: 2026-10-30T23:09:57.760Z

• Model: gpt-6.1-sol

This PDF is a classified, derived reading view. Raw JSONL remains the canonical lossless representation when included.

Unclassified user-role record – 2026-10-30T23:08:02.907Z

AGENTS.md instructions for D:\Temp\Sonrisa

<INSTRUCTIONS>

Project guidance for Codex

This repository is a 24-hour-scale interview exercise. The evaluation prioritizes the **authentic, inspectable reasoning and validation process** over feature breadth.

Product boundaries

Read `docs/PRODUCT_SCOPE.md`. Build an ASP.NET Core MVC / Razor, EF Core + SQLite prototype for users to define event alert rules and view matching events; SMTP and Slack webhook delivery; admin monitoring. Use a demo/seed event source and a clean source adapter seam. Avoid speculative third-party API integrations.

Non-negotiable engineering rules

- Do not claim success without executing checks and reporting command, exit code and outcome.

- Never invent human decisions, review findings, test results, screenshots, prompts or timestamps. Mark unperformed checks as NOT RUN.

- Do not quietly expand scope. State assumptions and tradeoffs.

- A matching event should generate at most one delivery per rule/channel/event despite retries; demonstrate idempotency.

- Avoid embedding secrets and credentials in source or logs. Keep .env, databases and real credentials untracked.

- Separate input/event normalization, matching rules, dispatch and delivery transports.

- No external send in automated tests. Stub email and Slack transports.

- Reviewer phases are read-only; do not modify files or make commits when reviewing.

- Do not commit, push, or publish from Codex. The external human-gated workflow owns commits after approval.

Evidence

The workflow launcher preserves actual AI prompts and responses in `docs/ai-process/runs/`; document architecture decisions in `docs/decisions.md` and real limitations in `docs/known-limitations.md`. Report uncertainty. When a review finds an issue, reference the affected files and an exact reproduction/check if possible.

</INSTRUCTIONS>

<environment_context>
<cwd>D:\Temp\Sonrisa</cwd>
<shell>powershell</shell>
<current_date>2026-10-31</current_date>
<timezone>Europe/Budapest</timezone>
<filesystem><workspace_roots><root>D:\Temp\Sonrisa</root></workspace_roots><permission_profile type="managed"><file_system type="restricted"><entry access="read"><special>:root</special></entry></file_system></permission_profile></filesystem>
</environment_context>

Unclassified user-role record – 2026-10-30T23:08:03.082Z

You are an INDEPENDENT SKEPTICAL REVIEWER, NOT a builder. Read-only. Do not modify code or files or commit.
Milestone: 01-plan - Scope, assumptions, design and acceptance criteria
Expected acceptance: Concrete small MVP, explicit product definition, risks, architecture and measurable success criteria.
Builder summary (untrusted):
Completed **01-plan**, with documentation only.

Paths added or updated:

- README [local file not included]
- Scope, assumptions, risks and 14 acceptance criteria [local file not included]
- Architecture and data model [local file not included]
- UI/HTTP flows and fixtures [local file not included]
- Validation strategy [local file not included]
- Decisions [local file not included], limitations [local file not included], process notes [local file not included]
- Executed commands and results [local file not included]

Verification actually run:

- Python documentation checks: **exit 0**; 29 local links resolve, all 14 criteria have validation scenarios, and no application files exist.
- Hash comparison: **exit 0**; original logs and all non-target starter files unchanged.
- Corrected Git status/diff/ignore checks: **exit 0**.
- SDK inspection: **exit 0**; only SDK **7.0.202** installed.
- Failures recorded: initial Git ownership check **exit 1**, `NUL` exclude-file check **exit 128**, and a rejected patch-format attempt. All were corrected.

Assumptions include one local demo user, fixed categories, severity 1–5, forward-only matching and three delivery attempts. Job uniqueness does not guarantee duplicate-free external notifications.

Application build/tests, browser checks, live delivery and independent review: **NOT RUN**. Before milestone 02, a compatible .NET 10 SDK is needed. Product assumptions remain for the normal human gate.

No application code, commits, pushes or authentic-log edits.

Deterministic validation: NOT RUN — see docs\ai-process\runs\01-plan-validation-0.md.
Read current repository diff and relevant code/tests. Look for broken behavior, data leaks, missing access controls, duplicate delivery, broken exception handling, unrealistic assumptions and insufficient tests, as applicable. Challenge the builder's claims; verify via concrete file references. Distinguish confirmed defects from risks and unverified items. Never manufacture issues to make the log look interesting.
Respond with: VERDICT (PASS / NEEDS_FIX / BLOCKED); critical and major findings with file paths and reproduction suggestions; test gaps; which claims you actually verified; suggested minimal fixes. A PASS is allowed if nothing material is found.

Assistant – 2026-10-30T23:08:05.112Z

I'll review the repository diff and planning documents, then check the builder's claims against the recorded evidence. I'll keep the review read-only and distinguish defects from risks and unverified work.

Assistant – 2026-10-30T23:09:03.452Z

The repository has no commits, so `git diff` provides no baseline; all files are untracked. The documents disclose this limitation. I found no application or tests, which is expected for this milestone. I'm now checking the delivery design and acceptance criteria for contradictions, and comparing the recorded hashes with the current files.

Assistant – 2026-10-30T23:09:57.728Z

VERDICT: PASS – milestone 01-plan only.

No critical or major findings. The plan defines a bounded MVP, explicit assumptions, separated responsibilities, risks and 14 measurable acceptance criteria with corresponding validation scenarios. See plan.md [local file not included] and architecture.md [local file not included].

The delivery design correctly distinguishes database job uniqueness from external effects. SMTP/Slack notifications may duplicate after an ambiguous timeout or crash; this limitation is explicit and consistent with PRODUCT_SCOPE. See architecture.md:91 [local file not included].

Checks I actually executed:

Command/check	Exit	Outcome
Read-only Python documentation/hash checks	0	Nine documents pass hygiene checks; 29 local links resolve; AC-01–14 align. Against the recorded baseline, 13 original files match and only the two declared placeholders differ. Seven declared documents plus four subsequent launcher artifacts exist.
`dotnet --list-sdks`	0	Only SDK 7.0.202 listed.
Scoped `git … status --short --untracked-files=all`	0	All files untracked.
Scoped `git … diff --check`	0	Empty; provides no content validation because files are untracked.
Scoped `git … check-ignore`	0	Representative environment, database, WAL and development configuration paths ignored.
`git … log -5 --oneline`	1	No commits; no Git baseline available.

The hash comparison corroborates the **recorded** baseline; it is not an independently preserved pre-builder snapshot.

Test gaps: application build/tests, SQLite concurrency and constraints, recovery, access controls, encoding, transport failures, browser behavior and live delivery are **NOT RUN**. No application exists, as required for this milestone. Their planned checks are documented in validation-strategy.md [local file not included]; runtime correctness remains unverified.

Minimal fixes: none required for 01-plan. Before 02-core, provide a compatible SDK and execute/pin the proposed setup. The .NET 10 LTS choice is supported by Microsoft's policy (https://dotnet.microsoft.com/en-us/platform/support/policy).

No files modified or commits made.