Codex session 01a18f06…

• Project: Sonrisa

• Storage: active

• Session ID: 01a18f06-5c92-74c2-8949-c253db7a8ab4

• Started: 2026-10-30T22:55:22.319Z

• Updated: 2026-10-30T23:07:59.732Z

• Model: gpt-6.1-sol

This PDF is a classified, derived reading view. Raw JSONL remains the canonical lossless representation when included.

Unclassified user-role record – 2026-10-30T22:55:23.852Z

AGENTS.md instructions for D:\Temp\Sonrisa

<INSTRUCTIONS>

Project guidance for Codex

This repository is a 24-hour-scale interview exercise. The evaluation prioritizes the **authentic, inspectable reasoning and validation process** over feature breadth.

Product boundaries

Read `docs/PRODUCT_SCOPE.md`. Build an ASP.NET Core MVC / Razor, EF Core + SQLite prototype for users to define event alert rules and view matching events; SMTP and Slack webhook delivery; admin monitoring. Use a demo/seed event source and a clean source adapter seam. Avoid speculative third-party API integrations.

Non-negotiable engineering rules

- Do not claim success without executing checks and reporting command, exit code and outcome.

- Never invent human decisions, review findings, test results, screenshots, prompts or timestamps. Mark unperformed checks as NOT RUN.

- Do not quietly expand scope. State assumptions and tradeoffs.

- A matching event should generate at most one delivery per rule/channel/event despite retries; demonstrate idempotency.

- Avoid embedding secrets and credentials in source or logs. Keep .env, databases and real credentials untracked.

- Separate input/event normalization, matching rules, dispatch and delivery transports.

- No external send in automated tests. Stub email and Slack transports.

- Reviewer phases are read-only; do not modify files or make commits when reviewing.

- Do not commit, push, or publish from Codex. The external human-gated workflow owns commits after approval.

Evidence

The workflow launcher preserves actual AI prompts and responses in `docs/ai-process/runs/`; document architecture decisions in `docs/decisions.md` and real limitations in `docs/known-limitations.md`. Report uncertainty. When a review finds an issue, reference the affected files and an exact reproduction/check if possible.

</INSTRUCTIONS>

<environment_context>
<cwd>D:\Temp\Sonrisa</cwd>
<shell>powershell</shell>
<current_date>2026-10-30</current_date>
<timezone>Europe/Budapest</timezone>
<filesystem><workspace_roots><root>D:\Temp\Sonrisa</root></workspace_roots><permission_profile type="managed"><file_system type="restricted"><entry access="read"><special>:root</special></entry><entry access="write"><path>D:\Temp\Sonrisa</path></entry><entry access="write"><special>:slash_tmp</special></entry><entry access="write"><special>:tmpdir</special></entry><entry access="read"><path>D:\Temp\Sonrisa\.git</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.agents</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.codex</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.aws</path></entry></file_system></permission_profile></filesystem>
</environment_context>

Unclassified user-role record – 2026-10-30T22:55:23.953Z

You are the BUILDER on milestone 01-plan: Scope, assumptions, design and acceptance criteria. Read AGENTS.md, docs/PRODUCT_SCOPE.md, workflow/milestones.json and current project files. Implement only the following goal: Write initial README, scope/assumptions, architecture, data model, API/UI flows, prioritized acceptance criteria, validation strategy, decision log and explicit exclusions. Do not implement application code yet.
Acceptance: Concrete small MVP, explicit product definition, risks, architecture and measurable success criteria.
Human constraints for this milestone:
(none)

Use modest, coherent engineering and realistic tests. Keep running verification while you work. Do not modify existing authentic AI log files. NEVER git commit or push. Report what you changed, paths touched, commands/tests actually run (including failures), assumptions, known limitations and anything needing human judgment. Mark not-tested things explicitly. Do not fake

review findings or evidence. If the scope is ambiguous, make a reasonable documented assumption rather than fabricate a requirement.

Assistant – 2026-10-30T22:55:26.674Z

I'll read the project guidance and current files, then write the MVP plan and measurable acceptance criteria without adding application code. I'll verify document consistency and report the checks actually run.

Assistant – 2026-10-30T22:56:00.444Z

The repository contains workflow scaffolding but no application yet. I'll plan for one local demo user, deterministic events, and dry-run delivery by default. I'll also distinguish database deduplication from external delivery guarantees: retrying an uncertain SMTP or Slack result can still produce a duplicate notification.

The environment has only .NET SDK 7.0.202. Git status initially failed on an ownership check; a command-scoped exception allowed read-only inspection without changing Git configuration.

Assistant – 2026-10-30T22:57:51.036Z

I'm defining matching and replay behavior so later tests have a clear target: severity 1–5, exact category, optional case-insensitive keyword, and no rematching of duplicate events. Rule edits and disabling will affect new events; existing delivery jobs will retain their snapshots.

The plan will target .NET 10 LTS, which requires a newer SDK than this machine has. Microsoft lists it as supported in its .NET support policy (https://dotnet.microsoft.com/en-us/platform/support/policy).

Unclassified user-role record – 2026-10-30T23:02:13.155Z

<environment_context>
  <current_date>2026-10-31</current_date>
  <timezone>Europe/Budapest</timezone>
  <filesystem><workspace_roots><root>D:\Temp\Sonrisa</root></workspace_roots><permission_profile type="managed"><file_system type="restricted"><entry access="read"><special>:root</special></entry><entry access="write"><path>D:\Temp\Sonrisa</path></entry><entry access="write"><special>:slash_tmp</special></entry><entry access="write"><special>:tmpdir</special></entry><entry access="read"><path>D:\Temp\Sonrisa\.git</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.agents</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.codex</path></entry><entry access="read"><path>D:\Temp\Sonrisa\.aws</path></entry></file_system></permission_profile></filesystem>
</environment_context>

Assistant – 2026-10-30T23:02:22.897Z

The document content is ready for checks, but the first write attempt was rejected because I used an invalid patch format. I'm correcting that and will verify the files afterward. The write failure will be recorded alongside the checks; application validation remains NOT RUN.

Assistant – 2026-10-30T23:03:31.177Z

The README, plan and architecture are now written. They separate event ingestion, matching, dispatch and transports, and define database constraints for deduplication. The acceptance criteria include replay, concurrency, rollback and retry checks.

Next I'm adding the HTTP/UI contracts and validation scenarios, then checking links, criterion coverage and the file-change boundary.

Assistant – 2026-10-30T23:04:56.323Z

The plan now maps 14 acceptance criteria to proposed checks, including the expected event, match and job counts for the demo. It also excludes provider integrations, backfill, resend controls and a public REST API.

I'm running documentation checks and confirming that the original workflow files and authentic logs remain unchanged.

Assistant – 2026-10-30T23:07:00.283Z

The documentation check passed: 14 criteria have matching validation scenarios, 29 local links resolve, and no application files were added. The hash comparison confirms both original launcher logs and all non-target starter files are unchanged.

A Git check also failed with `core.excludesFile=NUL`; using an empty command-scoped value fixed it. I'm recording both results and the remaining .NET SDK prerequisite in the evidence file.

Assistant – 2026-10-30T23:07:59.701Z

Completed **01-plan**, with documentation only.

Paths added or updated:

- README [local file not included]
- Scope, assumptions, risks and 14 acceptance criteria [local file not included]
- Architecture and data model [local file not included]
- UI/HTTP flows and fixtures [local file not included]
- Validation strategy [local file not included]
- Decisions [local file not included], limitations [local file not included], process notes [local

file not included]

- Executed commands and results [local file not included]

Verification actually run:

- Python documentation checks: **exit 0**; 29 local links resolve, all 14 criteria have validation scenarios, and no application files exist.
- Hash comparison: **exit 0**; original logs and all non-target starter files unchanged.
- Corrected Git status/diff/ignore checks: **exit 0**.
- SDK inspection: **exit 0**; only SDK **7.0.202** installed.
- Failures recorded: initial Git ownership check **exit 1**, `NUL` exclude-file check **exit 128**, and a rejected patch-format attempt. All were corrected.

Assumptions include one local demo user, fixed categories, severity 1–5, forward-only matching and three delivery attempts. Job uniqueness does not guarantee duplicate-free external notifications.

Application build/tests, browser checks, live delivery and independent review: **NOT RUN**. Before milestone 02, a compatible .NET 10 SDK is needed. Product assumptions remain for the normal human gate.

No application code, commits, pushes or authentic-log edits.